

# 인공지능

22년 삼성 AI 전문가과정  
6월 7일 화요일 3교시  
장병탁



## 3차시 : Problem-Solving Agents

서울대학교 컴퓨터공학부  
담당 교수: 장병탁

Seoul National University  
Byoung-Tak Zhang



# Lecture Overview

## 인공지능

### 3차시 : Problem-Solving Agents

서울대학교 컴퓨터공학부  
담당 교수: 장병탁


Seoul National University  
Byoung-Tak Zhang



exhaustive search,

## Lecture 3. Problem-Solving Agents

- Problem-Solving Agents
  - What are the **Problem-Solving Agents**?
  - Examples of Problems, Goals, Elements
- **Searching for Solutions** == solving problem
  - **Tree search** algorithms
- Uninformed Search Strategies
  - **Breadth-first search, Uniform-cost search, Depth-first search, Depth-limited search, Iterative-deepening search**
- Informed Search Strategies
  - **Greedy Search, A\* Search**
- Heuristic Functions
  - The 8-puzzle, **Relaxed problems**



7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Example:  
A typical instance of the 8-puzzle

# Problem-Solving Agents

## Reflex agents

## Goal-based agents

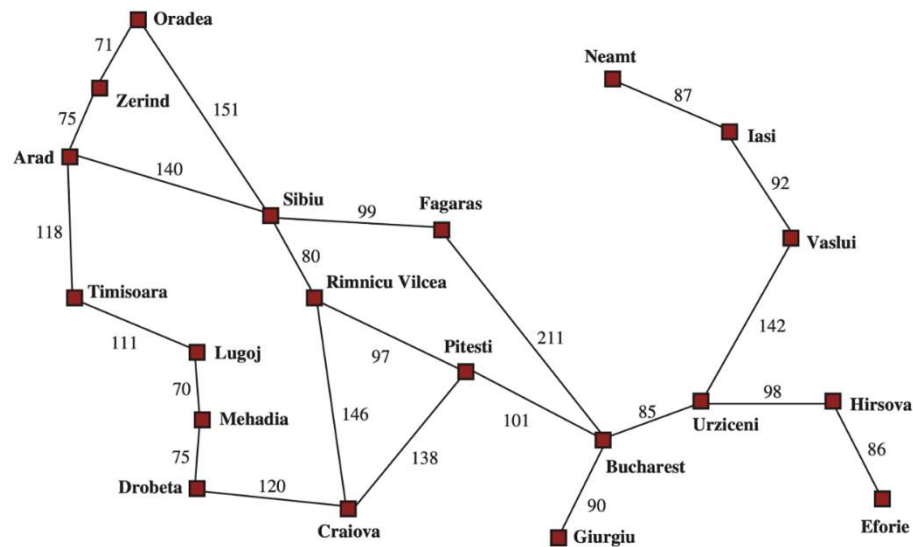
- Consider future actions and the desirability of their outcomes
- Problem-solving agents (Chapter 3)
  - Atomic representation (of states)
- Planning agents (Chapter 11)
  - Factored or structured representation

## Problem-solving agents (**this lecture**)

- Problems, solutions, search algorithms
- Uninformed search algorithms
- Informed search algorithms

# Problem Solving as Search in State Space: **Touring (Driving)**

Start State



Goal State

Goal: be in Bucharest

Problem formulation:


- states = various cities
- actions = drive between cities

Goal-based Agent

## Touring (Driving)

# Problem-Solving Agent: Vacuum Cleaner Robot

Start State



Goal State(s)

Goal: all rooms clean

Problem formulation:

- States: Boolean dirt and robot location
- Actions: Left, Right, Suck

## Vacuumping



# Problem Solving as Search in State Space

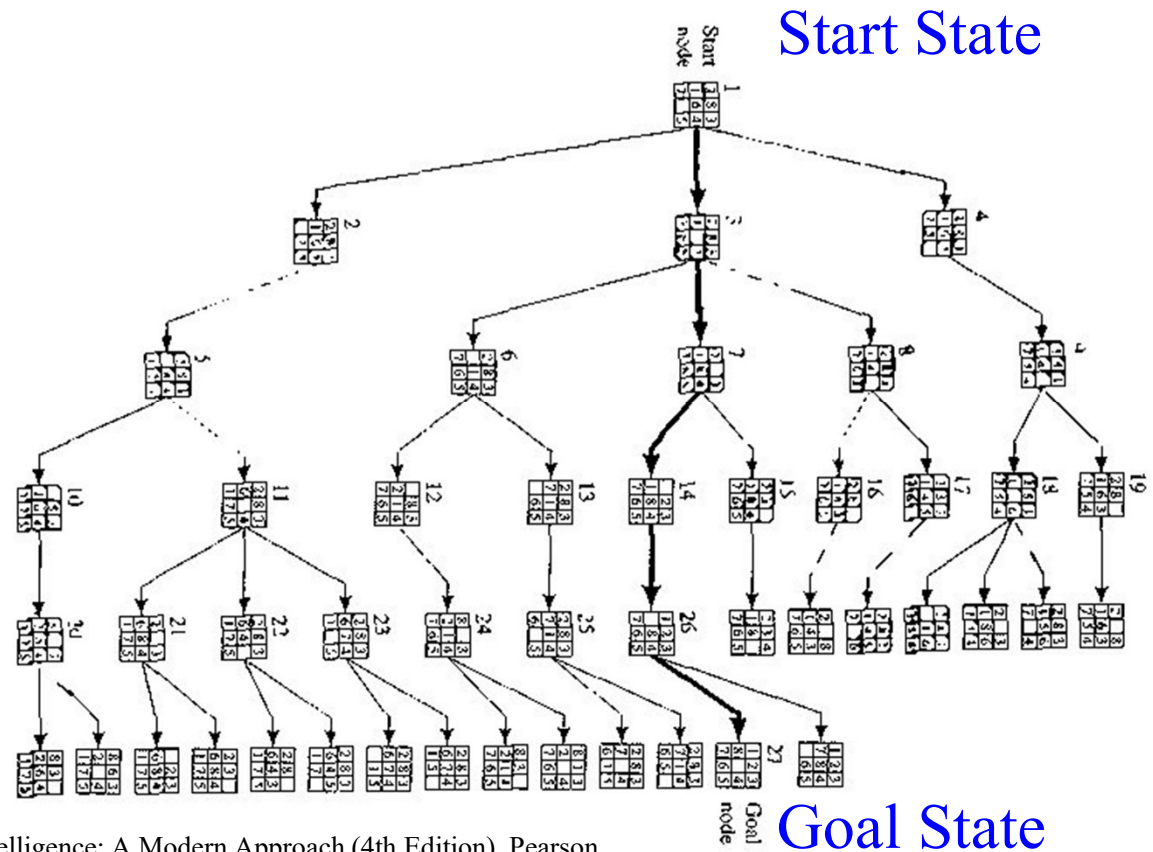
7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8


Goal State

Puzzle Solving



# Lecture 3. Problem-Solving Agents

- Problem-Solving Agents
  - What are the **Problem-Solving Agents**?
  - Examples of Problems, Goals, Elements
- Searching for Solutions
  - **Tree search** algorithms
- Uninformed Search Strategies
  - **Breadth-first search, Uniform-cost search, Depth-first search, Depth-limited search, Iterative-deepening search**
- Informed Search Strategies
  - **Greedy Search, A\* Search**
- Heuristic Functions
  - The 8-puzzle, **Relaxed problems**



7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Example:  
A typical instance of the 8-puzzle



## Outline (Lecture 3)

3.1 Problem-Solving Agents .....	10
3.2 Example Problems .....	15
3.3 Searching for Solutions .....	20
3.4 Uninformed Search Strategies .....	25
3.5 Informed Search Strategies .....	37
3.6 Heuristic Functions .....	45
Summary .....	49



## 3.1. Problem-Solving Agents



## 3.1 Problem-Solving Agents (1/4)

### Here we consider the environment:

- **Observable:** The agent always knows the current state
- **Known:** The agent knows which states are reached by each action
- **Deterministic:** Each action has exactly one outcome
- In this environment, the solution to any problem is a fixed sequence of actions. The agent ignores its percepts when choosing an action because it knows in advance what they will be. The agent is an open-loop system (“eyes closed”).

### Problem solving as search in state space

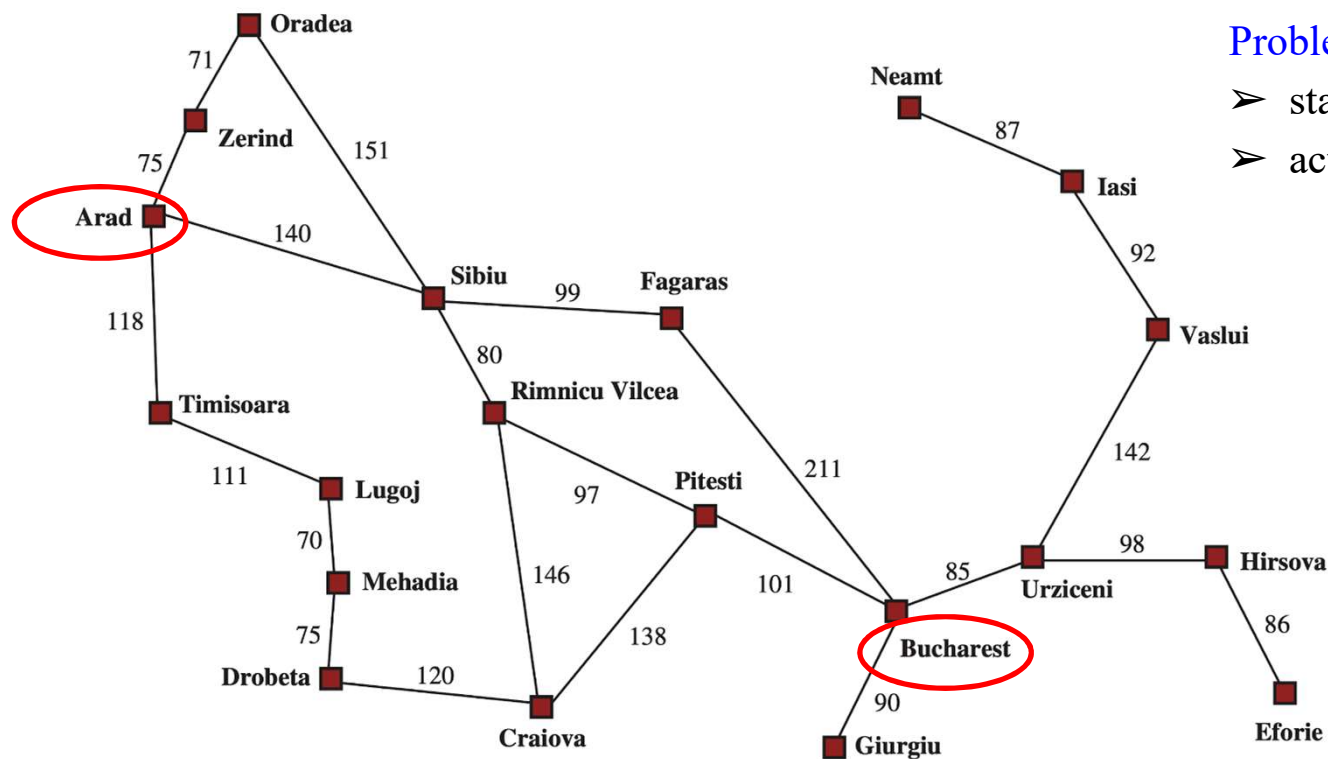
- **Goal formulation:** A goal is a set of world states to be reached.
- **Problem formulation:** The process of deciding what actions and states should be considered
- Search is the process of looking for a sequence of actions that reaches the goal in the state space.

### Well-defined problems and solutions

- **Problem** = <initial state, actions, transition model, goal test, path cost>
- A solution to a problem is an action sequence that leads from the initial state to a goal state.

## 3.1 Problem-Solving Agents (2/4)

Example: Romania



Goal: be in Bucharest

Problem formulation:

- states = various cities
- actions = drive between cities

## 3.1 Problem-Solving Agents (3/4)

Simple form of problem-solving agents

```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  persistent: seq, an action sequence, initially empty
               state, some description of the current world state
               goal, a goal, initially null
               problem, a problem formulation

  state ← UPDATE-STATE(state, percept)
  if seq is empty then
    goal ← FORMULATE-GOAL(state)
    problem ← FORMULATE-PROBLEM(state, goal)
    seq ← SEARCH(problem)
    if seq = failure then return a null action
  action ← FIRST(seq)
  seq ← REST(seq)
  return action
```

Percept  
State  
Action  
Goal

## 3.1 Problem-Solving Agents (4/4)

### Problem types

- Deterministic,  
Fully observable      ⇒ **Single-state**
- Non-observable      ⇒ **Conformant** problem  
(sensorless)
- Nondeterministic and/or  
Partially observable      ⇒ **Contingency** problem
- Unknown state space      ⇒ **Exploration** problem  
(online)


This lecture

Later






## 3.2. Example Problems



## 3.2 Example Problems (1/4)

Example: vacuum world

- **States:** Boolean dirt and robot location
- **Actions:** Left, Right, Suck
- **Transition model:** (next page)
- **Goal test:** no dirt
- **Path cost:** 1 per action



## 3.2 Example Problems (2/4)

### Vacuum world

➤ Single-state

start in 5

■ [R, S]

➤ Conformant

start in {1,2,3,4,5,6,7,8}


■ [R, S, L, S]

➤ Contingency

start in 5

Murphy's Law: [S] can make clean square dirty

■ [R, if dirt then S]



## 3.2 Example Problems (3/4)

Example: The 8-puzzle (NP-hard)

- **States:** integer location of tiles
- **Actions:** move blank tile Left, Right, Up, Down
- **Transition model:**
- **Goal test:** goal state
- **Path cost:** 1 per action

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

## 3.1 Problem-Solving Agents (4/4)


7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Puzzle Solving






## 3.3. Searching for Solutions





## 3.3 Searching for Solutions (1/4)

Partial search trees for  
finding a route from Arad  
to Bucharest



## 3.3 Searching for Solutions (2/4)

### Tree search algorithms

- **Basic Idea:** offline, simulated exploration of state space by generating successors of already-explored states


frontier

```
function TREE-SEARCH(problem) returns a solution, or failure
  initialize the frontier using the initial state of problem
  loop do
    if the frontier is empty then return failure
    choose a leaf node and remove it from the frontier
    if the node contains a goal state then return the corresponding solution
    expand the chosen node, adding the resulting nodes to the frontier
```

## 3.3 Searching for Solutions (3/4)

### Implementation: **states** vs **nodes**

- **State** is physical configuration
- Node is a data structure building block of a search tree
  - Parent, children, depth, path cost  $g(x)$



## 3.3 Searching for Solutions (4/4)


### Implementation: **EXPAND**

- The Expand function creates new nodes, filling in the various fields
- Use successor of the problem to create corresponding states

```
function EXPAND( node, problem ) returns a set of nodes
    successors ← the empty set
    for each action, result in SUCCESSOR-FN(problem, STATE[node]) do
        s ← a new NODE
        PARENT-NODE[s] ← node; ACTION[s] ← action; STATE[s] ← result
        PATH-COST[s] ← PATH-COST[node] + STEP-COST(node, action, s)
        DEPTH[s] ← DEPTH[node] + 1
        add s to successors
    return successors
```



## 3.4. Uninformed Search Strategies



## 3.4 Uninformed Search Strategies (1/11)

### Evaluating for search strategies

- Strategy is defined by picking the **order of node expansion**
- Strategies are evaluated along the following metrics:
  - **Completeness**: Guarantee to find a solution if one exists
  - **Optimality**: Guarantee to find a minimum-cost path
  - **Time complexity**: How many nodes to expand to find a solution
  - **Space complexity**: How many nodes to keep to perform search



## 3.4 Uninformed Search Strategies (2/11)

### 2) Uninformed search strategies

➤ Uninformed strategies use only the information available in the problem definition

- Breadth-first search
- Uniform-cost search
- Depth-first search
- Depth-limited search
- Iterative-deepening search

## 3.4 Uninformed Search Strategies (3/11)

### I. Breadth-first search

- Expand **shallowest unexpanded node**
- Implement: fringe is a **FIFO queue**, i.e., new successors go at end

**function** BREADTH-FIRST-SEARCH(*problem*) **returns** a solution, or failure

*node* ← a node with STATE = *problem*.INITIAL-STATE, PATH-COST = 0

**if** *problem*.GOAL-TEST(*node*.STATE) **then return** SOLUTION(*node*)

*frontier* ← a FIFO queue with *node* as the only element

*explored* ← an empty set

**loop do**

**if** EMPTY?(*frontier*) **then return** failure

*node* ← POP(*frontier*) /\* chooses the shallowest node in *frontier* \*/

    add *node*.STATE to *explored*

**for each** *action* **in** *problem*.ACTIONS(*node*.STATE) **do**

*child* ← CHILD-NODE(*problem*, *node*, *action*)


**if** *child*.STATE is not in *explored* or *frontier* **then**

**if** *problem*.GOAL-TEST(*child*.STATE) **then return** SOLUTION(*child*)

*frontier* ← INSERT(*child*, *frontier*)

## 3.4 Uninformed Search Strategies (4/11)

### I. Breadth-first search



<출처> Stuart J. Russell and Peter Norvig (2021). Artificial Intelligence: A Modern Approach (4th Edition). Pearson

- Complete? **Yes**
- Time?  $1 + b + b^2 + b^3 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$
- Space?  $O(b^{d+1})$  (keeps every node in memory)
- Optimal? **Yes**, Still, space is the big problem!

## 3.4 Uninformed Search Strategies (5/11)

### II. Uniform-cost search

- Expand **least-cost** unexpanded node
- Implementation: queue ordered by path cost  $g(n)$ , lowest first
- Complete? Yes, if step cost  $\geq \epsilon$
- Time? # of nodes with  $g \leq$  cost of optimal solution  $O(b^{\lceil C^*/\epsilon \rceil})$
- Space? # of nodes with  $g \leq$  cost of optimal solution  $O(b^{\lceil C^*/\epsilon \rceil})$
- Optimal? Yes, node expanded in order of  $g(n)$

$C^*$  is cost of the optimal solution


## 3.4 Uninformed Search Strategies (6/11)

### III. Depth-first search (DFS)

- Expand **deepest node** (depth)
- Implementation: **LIFO queue**, i.e., put successors at front
- Complete? No, fails in infinite-depth space with loops
  - Complete in finite space
- Time?  $O(b^m)$ , bad if  $m$  is bigger than  $d$
- Space?  $O(bm)$
- Optimal? **No**

### 3.4 Uninformed Search Strategies (7/11)

# Depth-first search



32 / 49



## 3.4 Uninformed Search Strategies (8/11)

### IV. Depth-limited search (DLS)

➤ Depth-first search with **depth limit  $l$**

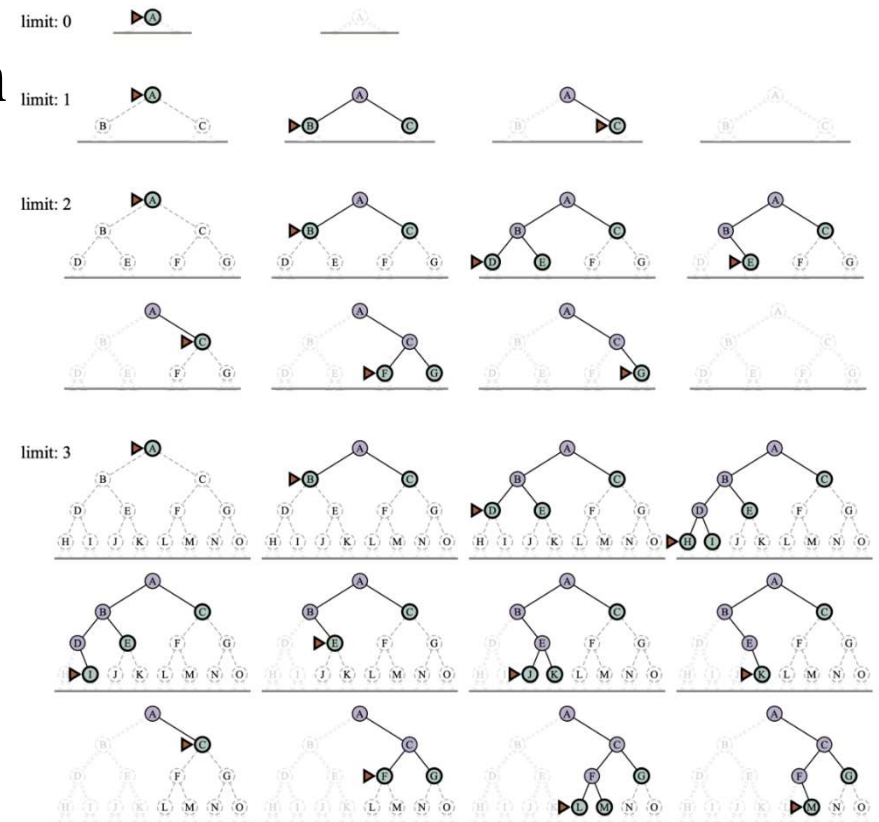
■ Node at depth  $l$ , have no successor

```
function DEPTH-LIMITED-SEARCH(problem,  $\ell$ ) returns a node or failure or cutoff  
  frontier  $\leftarrow$  a LIFO queue (stack) with NODE(problem.INITIAL) as an element  
  result  $\leftarrow$  failure  
  while not IS-EMPTY(frontier) do  
    node  $\leftarrow$  POP(frontier)  
    if problem.IS-GOAL(node.STATE) then return node  
    if DEPTH(node) >  $\ell$  then  
      result  $\leftarrow$  cutoff  
    else if not IS-CYCLE(node) do  
      for each child in EXPAND(problem, node) do  
        add child to frontier  
  return result
```

## 3.4 Uninformed Search Strategies (9/11)

### V. Iterative deepening search

- General strategy often used in combination with DFS and find the best depth limit



## 3.4 Uninformed Search Strategies (10/11)

### V. Iterative deepening search

- Complete? **Yes**
- Time?  $(d + 1)b^0 + db^1 + (d - 1)b^2 + \dots + b^d = O(b^d)$
- Space?  $O(bd)$
- Optimal? **Yes**, if step cost = 1

**function** ITERATIVE-DEEPENING-SEARCH(*problem*) **returns** a solution node or *failure*  
  **for** *depth* = 0 **to**  $\infty$  **do**  
    *result*  $\leftarrow$  DEPTH-LIMITED-SEARCH(*problem*, *depth*)  
    **if** *result*  $\neq$  *cutoff* **then return** *result*

## 3.4 Uninformed Search Strategies (11/11)

### Summary of algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes <sup>1</sup>	Yes <sup>1,2</sup>	No	No	Yes <sup>1</sup>	Yes <sup>1,4</sup>
Optimal cost?	Yes <sup>3</sup>	Yes	No	No	Yes <sup>3</sup>	Yes <sup>3,4</sup>
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$



## 3.5. Informed Search Strategies



## 3.5 Informed Search Strategies (1/7)

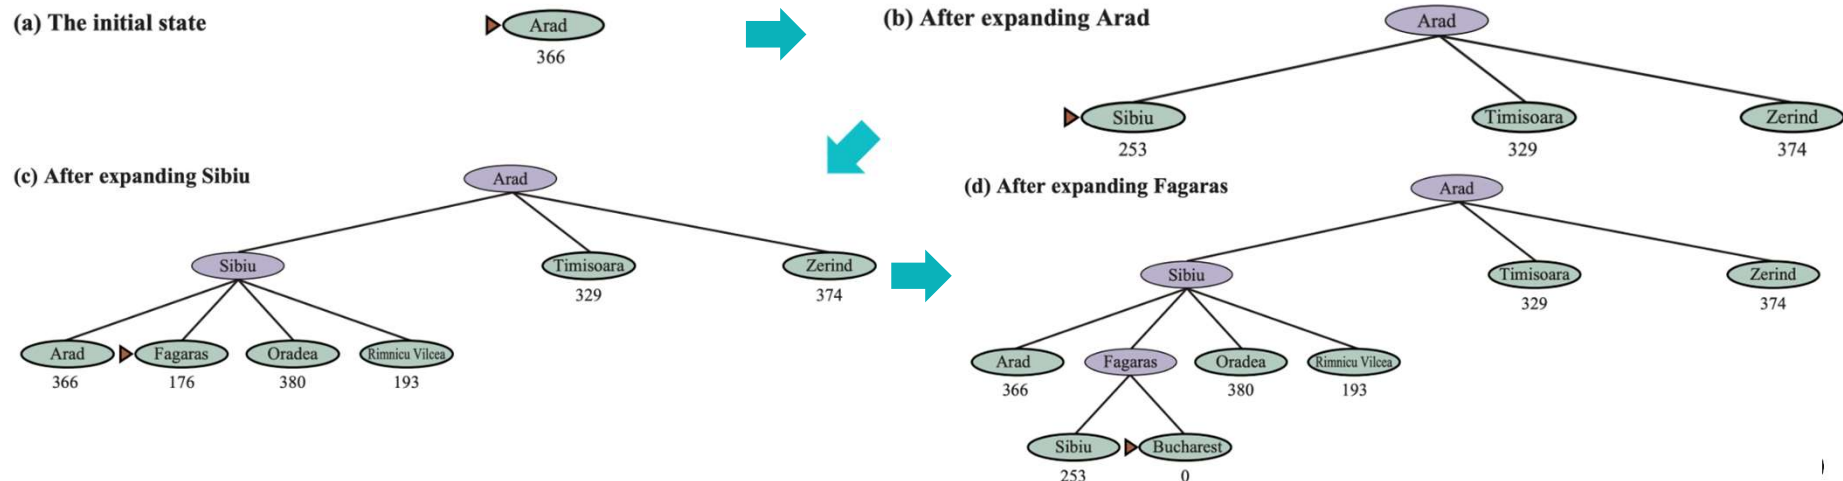
### Best-first search

- Idea: use an **evaluation function** for each node
  - Expand **more desirable unexpanded node**
- Implement: **queue sorted** in decreasing order of desirability
- Special cases:
  - **Greedy search**
  - **A\* search**

## 3.5 Informed Search Strategies (2/7)

### Greedy search

- Evaluation function  $h(n)$ 
  - Estimate of cost from  $n$  to the closest goal
- Greedy search expands the node that appears to be closest to goal



<출처> Stuart J. Russell and Peter Norvig (2021). Artificial Intelligence: A Modern Approach (4th Edition). Pearson



## 3.5 Informed Search Strategies (3/7)

### Greedy search

- Complete? No—can get stuck in loops,  
e.g., Iasi → Neamt → Iasi → Neamt → ...
- Complete? in finite space with repeated-state checking
- Time?  $O(b^m)$ , but a good heuristic can give dramatic improvement
- Space?  $O(b^m)$ , keeps all nodes in memory
- Optimal? No

## 3.5 Informed Search Strategies (4/7)

### A\* search

- Idea: avoid expanding paths that are already expensive
- Evaluation function  $f(n) = g(n) + h(n)$ 
  - $g(n)$  = cost so far to reach  $n$
  - $h(n)$  = estimated cost to goal from  $n$
  - $f(n)$  = estimated total cost of path through  $n$  to goal
- A\* search uses an admissible heuristic  $h(n) \leq h^*(n)$



## 3.5 Informed Search Strategies (5/7)

### A\* search

- Complete? Yes, unless there are infinitely many nodes with  $f \leq f(G)$
- Time? Exponential in [relative error in  $h \times$  length of soln.]
- Space? Keeps all nodes in memory
- Optimal? Yes—cannot expand  $f_{i+1}$  until  $f_i$  is finished

## 3.5 Informed Search Strategies (6/7)

### Conditions for optimality: Admissibility and consistency

- The first condition for optimality is that  $h(n)$  be an **admissible heuristic**, which *never overestimates* the cost to reach the goal.
  - $h(n) \leq h^*(n)$
  - Admissible heuristics are by nature *optimistic* because they think the cost of solving the problem is less than it actually is.
- A second, slightly stronger condition is **consistency** (or **monotonicity**) is required for applications of A\* to graph search.
  - $h(n) \leq c(n, a, n') + h(n')$
  - A heuristic  $h(n)$  is consistent if the estimated cost of reaching the goal from  $n$  is no greater than the step cost of getting to  $n'$  plus the estimated cost of reaching the goal from  $n'$ .

## 3.5 Informed Search Strategies (7/7)

### Optimality of A\*

$$h(n) \leq c(n, a, n') + h(n')$$

- A\* has the following properties:
  - The tree search version of A\* is optimal if  $h(n)$  is admissible, while the graph-search version is optimal if  $h(n)$  is consistent.
- We show the consistency of A\* in two steps.
- **Step 1:** If  $h(n)$  is consistent, then the value of  $f(n)$  along any path are nondecreasing.
  - $g(n') = g(n) + c(n, a, n')$
  - $f(n') = g(n') + h(n') = g(n) + c(n, a, n') + h(n') \geq g(n) + h(n) = f(n)$
- **Step 2:** Whenever A\* selects a node  $n$  for expansion, the optimal path to that node has been found.
  - Were this not the case, there would have to be another frontier node  $n'$  on the optimal path from the start node to  $n$ , (because  $f$  is nondecreasing along any path)  $n'$  would have lower  $f$ -cost than  $n$  and would have been selected first.



## 3.6. Heuristic Functions



## 3.6 Heuristic Functions (1/3)

### Example: the 8-puzzle

➤ Two heuristics:

$h_1(n)$  = number of misplaced tiles

$h_2(n)$  = total Manhattan distance

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

➤  $h_1(S) = 8$

➤  $h_2(n) = 3 + 1 + 2 + 2 + 2 + 3 + 3 + 2 = 18$



## 3.6 Heuristic Functions (2/3)

### Example: the 8-puzzle

➤ Dominance

■ If  $h_2(n) > h_1(n)$  for all  $n$ , then  $h_2$  **dominates**  $h_1$

➤ Typical search costs:

$d = 14$  IDS = 3,473,941 nodes

$A^*(h_1) = 539$  nodes

$A^*(h_2) = 113$  nodes

$d = 24$  IDS  $\approx$  54,000,000,000 nodes

$A^*(h_1) = 39,135$  nodes

$A^*(h_2) = 1,641$  nodes

➤ Given any admissible heuristics  $h_a, h_b$ ,

$h(n) = \max(h_a(n), h_b(n))$  is also admissible and dominates

## 3.6 Heuristic Functions (3/3)

### Example: the 8-puzzle

- Admissible heuristics can be derived from the exact solution cost of a relaxed version of the problem
- If a tile can move anywhere (in 8-puzzle)
  - then  $h_1(n)$  gives the shortest solution
- If a tile can move to any adjacent square (in 8-puzzle)
  - then  $h_2(n)$  gives the shortest solution
- Key point: the optimal solution cost of a relaxed problem is no greater than the optimal solution cost of the real problem
- Given any admissible heuristics  $h_a, h_b$ ,  
 $h(n) = \max(h_a(n), h_b(n))$  is also admissible and dominates  $h_a, h_b$

# Summary

1. **Heuristic functions** estimate **costs of shortest paths**
2. **Good heuristics can dramatically reduce search cost**
3. **Greedy best-first search** expands **lowest  $h$** 
  - incomplete and not always optimal
4. **A\* search** expands **lowest  $g + h$** 
  - Complete and optimal
  - Also optimally efficient (up to tie-breaks, for forward search)
5. **Admissible heuristics** can be derived from exact solution of relaxed problems