

인공지능

22년 삼성 AI 전문가과정
6월 7일 화요일 5교시
장병탁



6차시 : Logical Inference

서울대학교 컴퓨터공학부
담당 교수: 장병탁

Seoul National University
Byoung-Tak Zhang



Lecture Overview

인공지능

6차시 : Logical Inference

서울대학교 컴퓨터공학부
담당 교수: 장병탁

Seoul National University
Byoung-Tak Zhang



Introduction

- ❑ **Knowledge-based agents** (Previous lecture)
 - Represent knowledge about the world
 - Deduce the actions to take
- ❑ **Representing knowledge using logic** (Previous lecture)
 - Propositional logic (PL)
 - First-order logic (FOL)
- ❑ **Inference in first-order logic** (This lecture)
 - Propositionalization
 - Forward chaining
 - Backward chaining
 - Resolution inference

Knowledge in First-Order Logic: Colonel West Problem

Colonel West Problem

- The law says that it is a crime for an American to sell weapons to hostile nations. The country Nono, an enemy of America, has some missiles, and all of its missiles were sold to it by Colonel West, who is American.
- **Prove that Colonel West is a criminal.**

First-order logic description of the problem

$American(x) \wedge Weapon(y) \wedge Sells(x,y,z) \wedge Hostile(z) \Rightarrow Criminal(x)$

$Owns(Nono, MI)$

$Missile(MI)$

$Missile(x) \wedge Owns(Nono, x) \Rightarrow Sells(West, x, Nono)$

$Missile(x) \Rightarrow Weapon(x)$

$Enemy(x, America) \Rightarrow Hostile(x)$

$American(West)$

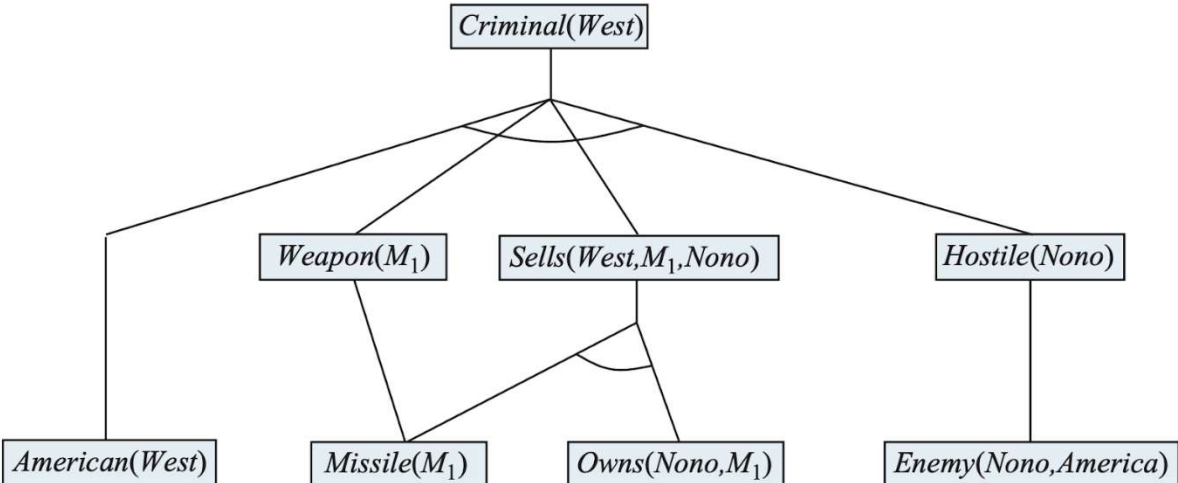
$Enemy(Nono, America)$

 $Criminal(West)?$

Inference in FOL: Proof Tree

- How can answer any answerable **first-order logic** question?
 - There are four major ways to make inferences in FOL.
 - Propositionalization
 - Forward chaining
 - Backward chaining
 - Resolution inference
- Proving if Colonel West is a criminal**
-
- ```
graph TD; A[Criminal(West)] --- B[];
```

## Proving if Colonel West is a criminal



출처: Stuart J. Russell and Peter Norvig (2021). Artificial Intelligence: A Modern Approach (4rd Edition). Pearson

## Outline (Lecture 6)

|                                                   |    |
|---------------------------------------------------|----|
| 6.1 Propositional vs. First-Order Inference ..... | 6  |
| 6.2 Unification and Lifting .....                 | 12 |
| 6.3 Forward Chaining .....                        | 17 |
| 6.4 Backward Chaining .....                       | 24 |
| 6.5 Resolution .....                              | 29 |
| Summary .....                                     | 34 |





## 6.1 Propositional vs. First-Order Inference



## 6.1 Propositional vs. First-Order Inference (1/5)

### Inference rules

➤ **Modus Ponens**

$$\frac{\alpha \Rightarrow \beta, \alpha}{\beta}$$

whenever any sentences of the form  $\alpha \Rightarrow \beta$  and  $\alpha$  are given, then the sentence  $\beta$  can be inferred.

➤ **And-Elimination**

$$\frac{\alpha \wedge \beta}{\alpha}$$

➤ **Resolution**

$$(\alpha_1, \dots, \alpha_j, \dots, \alpha_n), (\alpha_1, \dots, \neg \alpha_j, \dots, \alpha_n) \models (\alpha_1, \dots, \alpha_{j-1}, \alpha_{j+1}, \dots, \alpha_n)$$

➤ **Universal instantiation (UI)**

➤ **Existential instantiation (EI)**

**Both for PL & FOL**

**new for FOL**



## 6.1 Propositional vs. First-Order Inference (2/5)

### Inference rules for quantifiers

$$\forall x \text{ Smart}(x)$$

#### 1) Universal instantiation (UI)

- We can infer any sentence  $\alpha$  obtained by substituting a **ground term**  $g$  for the variable  $v$ :  $\frac{\forall v \alpha}{\text{Subst}(\{v/g\}, \alpha)}$

- $\forall x \text{ King}(x) \wedge \text{Greedy}(x) \Rightarrow \text{Evil}(x)$

$$\text{King}(\text{John}) \wedge \text{Greedy}(\text{John}) \Rightarrow \text{Evil}(\text{John})$$

**yields**  $\text{King}(\text{Richard}) \wedge \text{Greedy}(\text{Richard}) \Rightarrow \text{Evil}(\text{Richard})$

$$\text{King}(\text{Father}(\text{John})) \wedge \text{Greedy}(\text{Father}(\text{John})) \Rightarrow \text{Evil}(\text{Father}(\text{John}))$$

## 6.1 Propositional vs. First-Order Inference (3/5)

### 2) Existential instantiation (EI)

- The variable  $v$  in the sentence  $\alpha$  is replaced by a single new constant symbol  $k$ :

$$\frac{\exists v \alpha}{\text{Subst}(\{v/k\}, \alpha)}$$

$$\boxed{\exists x \text{ Smart}(x)}$$

- From the sentence

$$\exists x \text{ Crown}(x) \wedge \text{OnHead}(x, \text{John})$$

we can infer the sentence

$$\text{Crown}(C_1) \wedge \text{OnHead}(C_1, \text{John})$$

provided  $C_1$  is a new constant symbol, called a **Skolem** constant.

- **Inferentially equivalent** (but **not logically equivalent**) in the sense that it is satisfiable exactly when the original KB is satisfiable.

## 6.1 Propositional vs. First-Order Inference (4/5)

### Reduction to Propositional Inference

- Suppose the KB contains just the following

$$\forall x \text{ King}(x) \wedge \text{Greedy}(x) \Rightarrow \text{Evil}(x)$$

*King(John)*

*Greedy(John)*

*Brother(Richard, John)*

- **Instantiating the universal sentence** in all possible ways,  $\{x/\text{John}\}$  and  $\{x/\text{Richard}\}$ , we have

$$\text{King}(\text{John}) \wedge \text{Greedy}(\text{John}) \Rightarrow \text{Evil}(\text{John})$$

$$\text{King}(\text{Richard}) \wedge \text{Greedy}(\text{Richard}) \Rightarrow \text{Evil}(\text{Richard})$$

- **The new KB is propositionalized:**

*King(John), Greedy(John), Evil(John), King(Richard), etc.*

## 6.1 Propositional vs. First-Order Inference (5/5)

### Technique of Propositionalization

- First-order inference via propositionalization is complete— that is, any entailed sentence can be proved.
- What happens when the sentence is not entailed? We cannot tell.
- The question of entailment for first-order logic is **semidecidable**—that is, algorithms exist that say yes to every entailed sentence, but no algorithm exists that also says no to every non-entailed sentence.
- Cf. The halting problem for Turing machines.



## 6.2 Unification and Lifting



## 6.2 Unification and Lifting (1/4)

### Motivating Example

$$\forall x \text{ King}(x) \wedge \text{Greedy}(x) \Rightarrow \text{Evil}(x)$$

$$\text{King}(\text{John})$$

$$\forall y \text{ Greedy}(y)$$

- Propositionalization approach may generate

$$\text{King}(\text{Richard}) \wedge \text{Greedy}(\text{Richard}) \Rightarrow \text{Evil}(\text{Richard})$$

which does not match the KB and, thus, *useless* for proving *Evil(John)*.

- We can get the inference immediately if we can find a *substitution*  $\theta$  such that *King(x)* and *Greedy(x)* match *King(John)* and *Greedy(John)*
- $\theta = \{x/\text{John}, y/\text{John}\}$  works

## 6.2 Unification and Lifting (2/4)

### Unification

- $Unify(p, q) = \theta$  where  $p\theta = q\theta$  ( $\theta$ : most general unifier (MGU))
- Lifted inference rules require finding substitutions that make different logical expressions look identical. This process is called unification.

| $p$              | $q$                   | $\theta$                         |
|------------------|-----------------------|----------------------------------|
| $Knows(John, x)$ | $Knows(John, Jane)$   | $\{x / Jane\}$                   |
| $Knows(John, x)$ | $Knows(y, Bill)$      | $\{x / OJ, y / Jane\}$           |
| $Knows(John, x)$ | $Knows(y, Mother(y))$ | $\{y / Jone, x / Mother(John)\}$ |
| $Knows(John, x)$ | $Knows(x, Elizabeth)$ | $fail$                           |



## 6.2 Unification and Lifting (3/4)

### Generalized Modus Ponens (GMP)

$$\frac{p'_1, p'_2, \dots, p'_n, (p_1 \wedge p_2 \wedge \dots \wedge p_n \Rightarrow q)}{\text{Subst}(\theta, q)}$$

*King(John)*  
 $\forall y \text{ Greedy}(y)$   
 $\forall x \text{ King}(x) \wedge \text{Greedy}(x) \Rightarrow \text{Evil}(x)$

where  $\text{Subst}(\theta, p'_i) = \text{Subst}(\theta, p_i)$  or  $p'_i\theta = p_i\theta$  for  $\forall i$ .

$p'_1$  is *King(John)*

$p'_2$  is *Greedy(y)*

$\theta$  is  $\{x/\text{John}, y/\text{John}\}$

$\text{Subst}(\theta, q) = q\theta$  is *Evil(John)*

$p_1$  is *King(x)*

$p_2$  is *Greedy(x)*

$q$  is *Evil(x)*

GMP is a **lifted** version of MP. It raises MP from ground (variable-free) propositional logic to first-order logic.

## 6.2 Unification and Lifting (4/4)

### Soundness of GMP

$p \models p\theta$  by UI

$p\theta = \text{Subst}(\theta, p)$

$p'_1, \dots, p'_n \models p'_1\theta, \dots, p'_n\theta$

$p_1 \wedge \dots \wedge p_n \Rightarrow q \models p_1\theta \wedge \dots \wedge p_n\theta \Rightarrow q\theta$

Now  $\theta$  in GMP is defined as  $p'_i\theta = p_i\theta$  for  $\forall i$ .

Thus, we have

$p'_1, \dots, p'_n, (p_1 \wedge \dots \wedge p_n \Rightarrow q) \models q\theta$

provided that  $p'_i\theta = p_i\theta$  for  $\forall i$



## 6.3 Forward Chaining



## 6.3 Forward Chaining (1/6)

Inference methods for direct manipulation of FOL (unlike propositionalization)

### Example: Colonel West Problem

- The law says that it is a crime for an American to sell weapons to hostile nations. The country Nono, an enemy of America, has some missiles, and all of its missiles were sold to it by Colonel West, who is American.
- Prove that Colonel West is a criminal.

$American(x) \wedge Weapon(y) \wedge Sells(x,y,z) \wedge Hostile(z) \Rightarrow Criminal(x)$

$Owns(Nono, M1)$

$Missile(M1)$

$Missile(x) \wedge Owns(Nono, x) \Rightarrow Sells(West, x, Nono)$

$Missile(x) \Rightarrow Weapon(x)$

$Enemy(x, America) \Rightarrow Hostile(x)$

$American(West)$

$Enemy(Nono, America)$

**First-order definite clauses**

**Datalog:** First-order definite clauses with no function symbols

## 6.3 Forward Chaining (2/6)

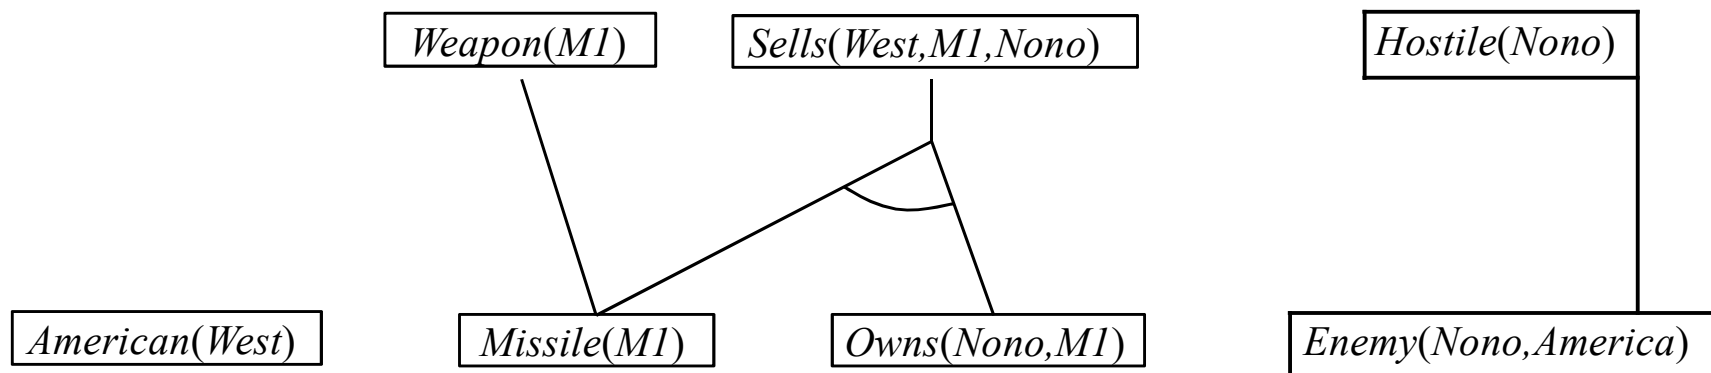
*American(West)*

*Missile(MI)*

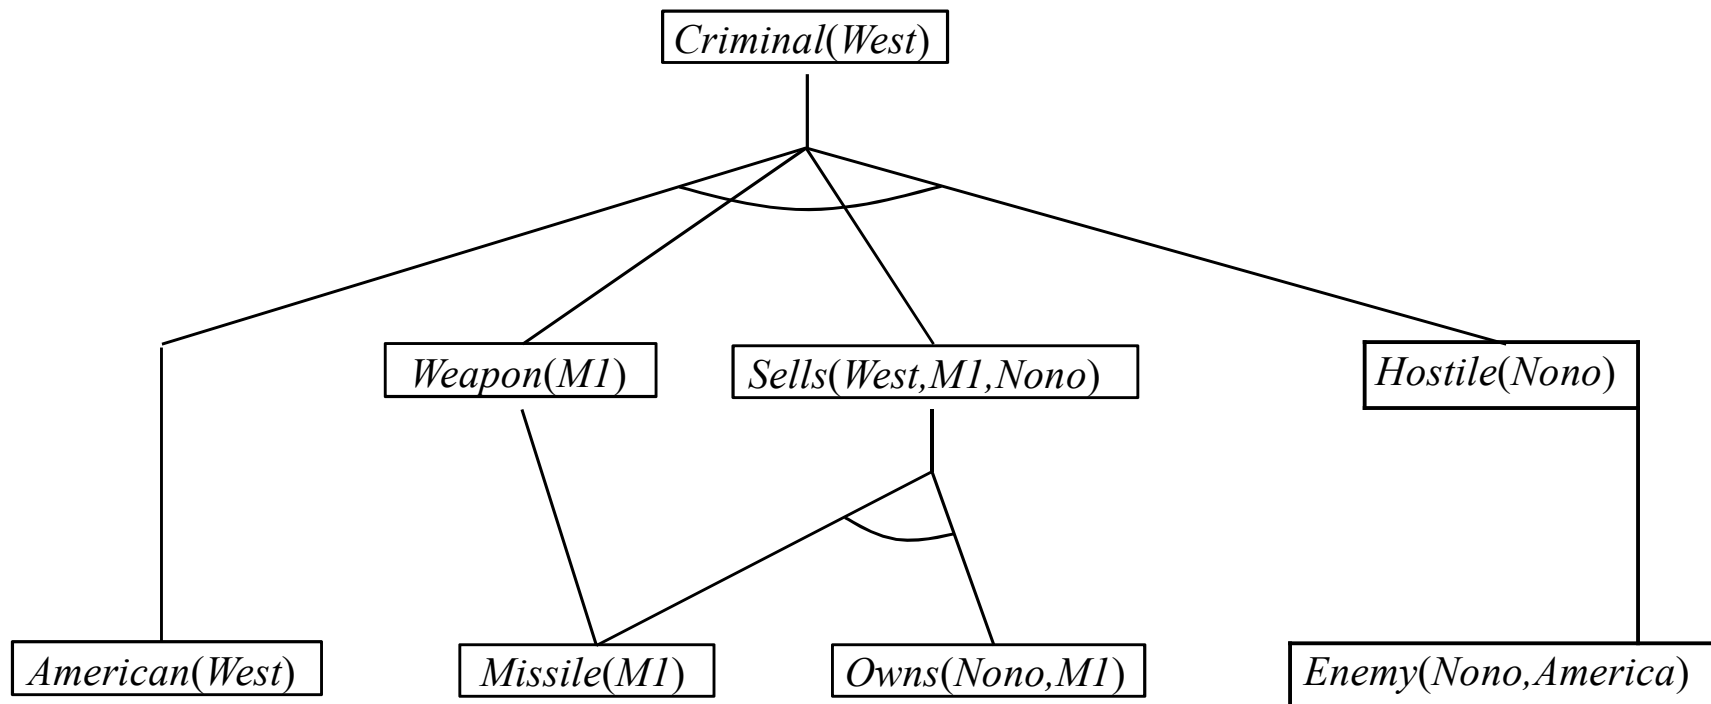
*Owns(Nono,MI)*

*Enemy(Nono,America)*

## 6.3 Forward Chaining (3/6)



## 6.3 Forward Chaining (4/6)





## 6.3 Forward Chaining (5/6)

### Forward chaining algorithm

```
function FOL-FC-ASK(KB, α) returns a substitution or false
 inputs: KB , the knowledge base, a set of first-order definite clauses
 α , the query, an atomic sentence

 while true do
 $new \leftarrow \{ \}$ // The set of new sentences inferred on each iteration
 for each rule in KB do
 $(p_1 \wedge \dots \wedge p_n \Rightarrow q) \leftarrow \text{STANDARDIZE-VARIABLES}(\text{rule})$
 for each θ such that $\text{SUBST}(\theta, p_1 \wedge \dots \wedge p_n) = \text{SUBST}(\theta, p'_1 \wedge \dots \wedge p'_n)$
 for some p'_1, \dots, p'_n in KB
 $q' \leftarrow \text{SUBST}(\theta, q)$
 if q' does not unify with some sentence already in KB or new then
 add q' to new
 $\phi \leftarrow \text{UNIFY}(q', \alpha)$
 if ϕ is not failure then return ϕ
 if $new = \{ \}$ then return false
 add new to KB
```

## 6.3 Forward Chaining (6/6)

### Properties of forward chaining

- Sound and complete for first-order definite clauses
  - Proof similar to propositional proof
- Datalog = first-order definite clauses + no functions (e.g., crime KB)
- FC terminates for Datalog in polynomial number of iterations
- May not terminate in general if  $\alpha$  is not entailed
- This is unavoidable: entailment with definite clauses is semidecidable
- Forward chaining is widely used in deductive databases



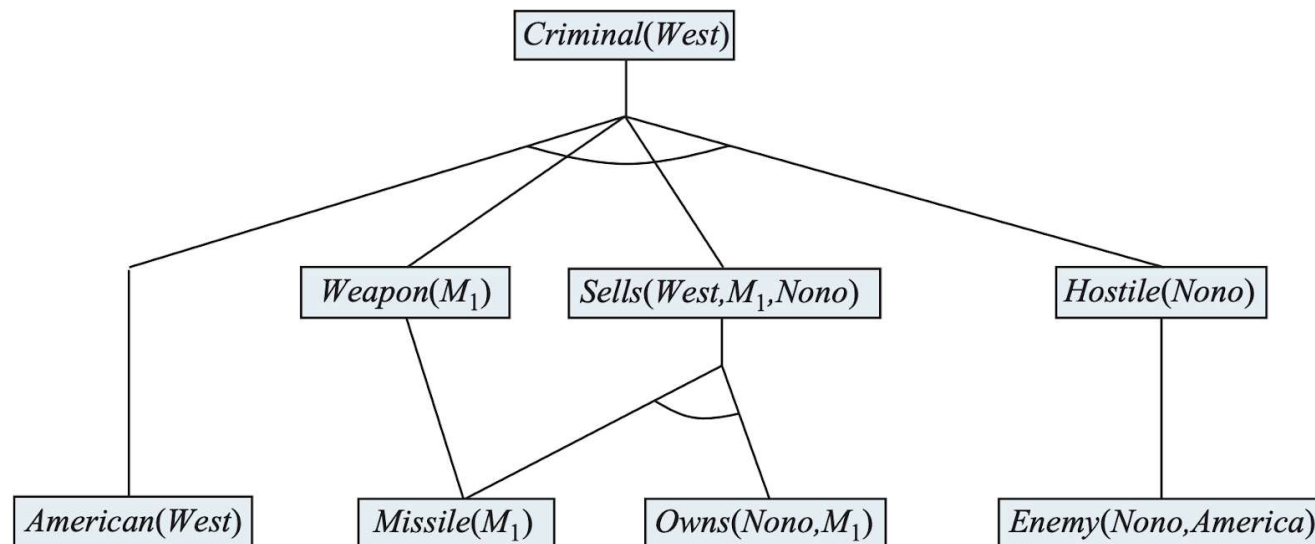
## 6.4 Backward Chaining



## 6.4 Backward Chaining (1/4)

**The proof tree generated by forward chaining on the crime example.**

- The initial facts appear at the **bottom level**, facts inferred on the first iteration in the **middle level**, and facts inferred on the second iteration at the **top level**.



## 6.4 Backward Chaining (2/4)

A simple **backward-chaining** algorithm for first-order knowledge bases

```
function FOL-BC-ASK(KB, query) returns a generator of substitutions
 return FOL-BC-OR(KB, query, { })
```

```
function FOL-BC-OR(KB, goal, θ) returns a substitution
 for each rule in FETCH-RULES-FOR-GOAL(KB, goal) do
 (lhs \Rightarrow rhs) \leftarrow STANDARDIZE-VARIABLES(rule)
 for each θ' in FOL-BC-AND(KB, lhs, UNIFY(rhs, goal, θ)) do
 yield θ'
```

```
function FOL-BC-AND(KB, goals, θ) returns a substitution
 if $\theta = \text{failure}$ then return
 else if LENGTH(goals) = 0 then yield θ
 else
 first, rest \leftarrow FIRST(goals), REST(goals)
 for each θ' in FOL-BC-OR(KB, SUBST(θ , first), θ) do
 for each θ'' in FOL-BC-AND(KB, rest, θ') do
 yield θ''
```

## 6.4 Backward Chaining (3/4)

### Properties of **backward chaining**

- **Depth-first** recursive proof search: space is linear in size of proof
- **Incomplete** due to infinite loops
  - fix by checking current goal against every goal on stack
- Inefficient due to repeated subgoals (both success and failure)
  - fix using caching of previous results (extra space!)
- Widely used (without improvements!) for **logic programming**

## 6.4 Backward Chaining (4/4)

### Logic programming (PROLOG)

- Algorithm = Logic + Control
- Prolog program = sets of definite clauses

`criminal(X) :- american(X), weapon(Y), sells(X, Y, Z), hostile(Z)`

`append([], Y, Y).`

`append([A|X], Y, [A|Z]) :- append(X, Y, Z)`

`append(X, Y, [1,2])`

`→`

`X=[]        Y=[1,2];`

`X=[1]      Y=[2];`

`X=[1,2]    Y=[]`

- Prolog uses database semantics, i.e. closed-world assumption and negation as failure
- Depth-first backward-chaining search





## 6.5 Resolution



## 6.5 Resolution (1/4)

### Resolution in full first-order

$$(A_1 \vee \dots \vee A_i \vee A_{i+1} \vee \dots \vee A_k, m_1 \vee \dots \vee m_j \vee m_{j+1} \vee \dots \vee m_n)\theta$$

where  $Unify(A_i, \neg m_j) = \theta$ .

For example,

$$\frac{\neg Rich(x) \vee Unhappy(x)}{Rich(Ken)} \quad Unhappy(Ken)$$

with  $\theta = \{x/Ken\}$

## 6.5 Resolution (2/4)

### Conversion to CNF

Everyone who loves all animals is loved by someone:

$$\forall x [\forall y \text{ Animal}(y) \Rightarrow \text{Loves}(x, y)] \Rightarrow [\exists y \text{ Loves}(y, x)]$$

#### 1. Eliminate biconditionals and implications

$$\forall x [\neg \forall y \neg \text{Animal}(y) \vee \text{Loves}(x, y)] \vee [\exists y \text{ Loves}(y, x)]$$

#### 2. Move $\neg$ inwards: $\neg \forall x p \equiv \exists x \neg p$ , $\neg \exists x p \equiv \forall x \neg p$ :

$$\forall x [\exists y \neg(\neg \text{Animal}(y) \vee \neg \text{Loves}(x, y))] \vee [\exists y \text{ Loves}(y, x)]$$

$$\forall x [\exists y \neg \neg \text{Animal}(y) \vee \neg \text{Loves}(x, y)] \vee [\exists y \text{ Loves}(y, x)]$$

$$\forall x [\exists y \text{ Animal}(y) \vee \neg \text{Loves}(x, y)] \vee [\exists y \text{ Loves}(y, x)]$$

## 6.5 Resolution (3/4)

### Conversion to CNF (contd.)

3. **Standardize variables**: each quantifier should use a different one

$$\forall x [\exists y \textit{Animal}(y) \wedge \neg \textit{Loves}(x, y)] \vee [\exists z \textit{Loves}(z, x)]$$

4. **Skolemize**: a more general form of existential instantiation. Each existential variable is replaced by a **Skolem function** of the enclosing universally quantified variables:

$$\forall x [\textit{Animal}(F(x)) \wedge \neg \textit{Loves}(x, F(x))] \vee \textit{Loves}(G(x), x)$$

5. **Drop universal quantifiers**:

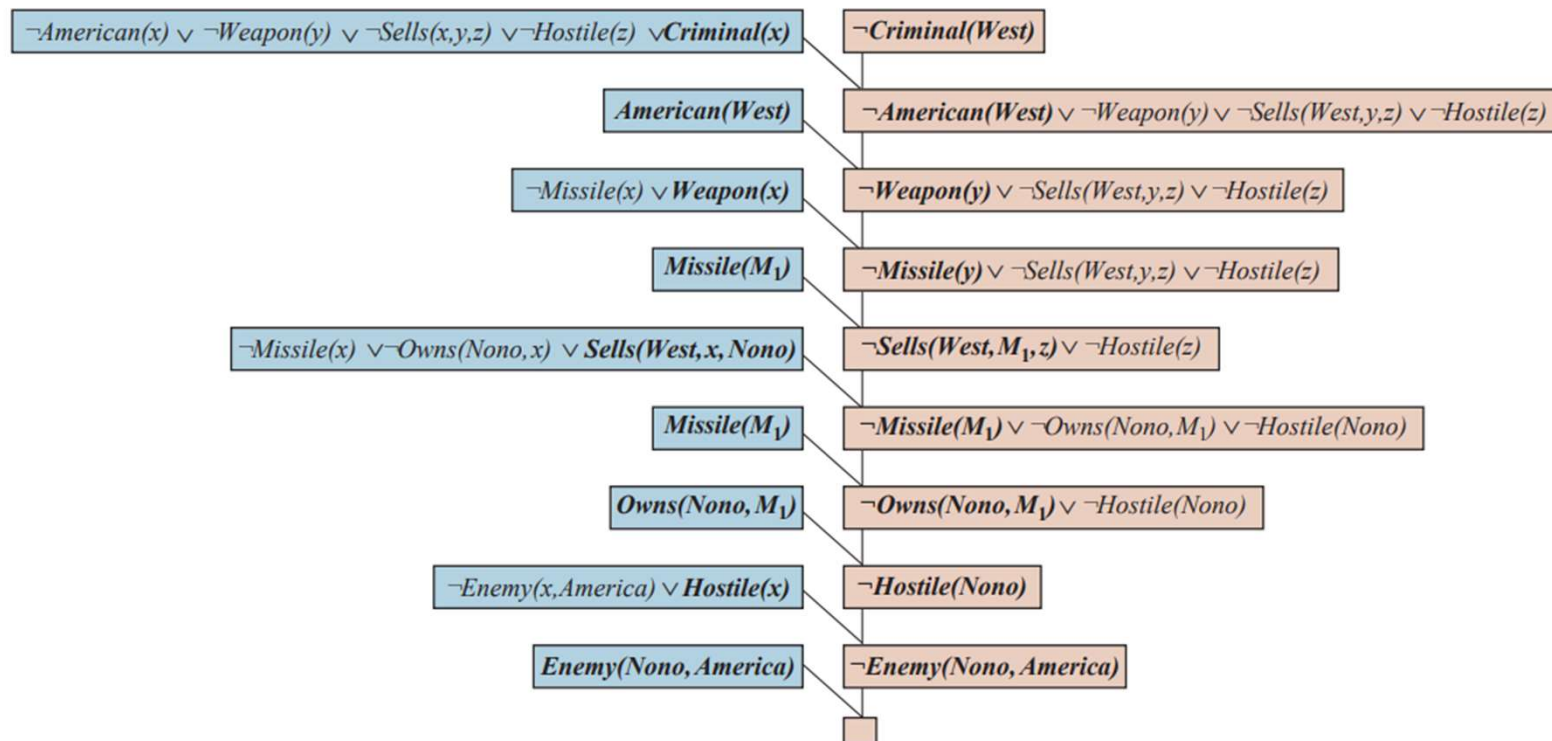
$$[\textit{Animal}(F(x)) \wedge \neg \textit{Loves}(x, F(x))] \vee \textit{Loves}(G(x), x)$$

6. **Distribute over**:

$$[\textit{Animal}(F(x)) \vee \textit{Loves}(G(x), x)] \wedge [\neg \textit{Loves}(x, F(x)) \vee \textit{Loves}(G(x), x)] \leftarrow \textbf{CNF}$$

## 6.5 Resolution (4/4)

### Example proofs



# Summary

- Instantiation is slow, unless the domain is small.
- The use of **unification** to identify appropriate substitutions for variables eliminates the instantiation step in FO proofs, making the process more efficient.
- A lifted version of **Modus Ponens** uses unification to provide a powerful inference rule, **generalized Modus Ponens**. The forward-chaining and backward-chaining algorithms apply this rule to sets of definite clauses.
- **Forward chaining** is used in deductive databases and production systems. Forward chaining is complete for Datalog.
- **Prolog**, unlike first-order logic, uses a closed world with the unique names assumption and negation as failure.
- The generalized **resolution** inference rule provides a complete proof system for first-order logic, using knowledge bases in CNF.