

인공지능

22년 삼성 AI 전문가과정
6월 9일 목요일 2교시
장병탁



17차시 : Reinforcement Learning

서울대학교 컴퓨터공학부
담당 교수: 장병탁

Seoul National University
Byoung-Tak Zhang



Lecture Overview

인공지능

17차시 : Reinforcement Learning

서울대학교 컴퓨터공학부
담당 교수: 장병탁

Seoul National University
Byoung-Tak Zhang



Introduction: Reinforcement Learning

❑ Supervised Learning (Previous lectures)

- With **supervised learning**, an agent learns by **passively observing example input/output pairs** provided by a “**teacher**.”
- **Deep learning** models also learn this way, including feedforward neural networks, convolutional neural networks, and recurrent neural networks.

❑ Reinforcement Learning (This lecture)

- In reinforcement learning, the agents can **actively learn from their own experience, without a teacher**, by considering their own ultimate success or failure.
- We see how experiencing **rewards** and **punishments** can teach an agent **how to maximize rewards** in the future
- Passive/Active RL, Generalization in RL, Policy Search, Inverse RL, etc.

Problems with Supervised Learning

What's wrong with supervised learning?


- ❑ A supervised learning agent learns by **passively** observing example input/output pairs provided by a “**teacher**”
 - ❑ Can we train the chess agent with supervised learning?
 - Data: examples of chess positions (each labeled with the correct move)
 - Available training data set for winners: 10^8
 - The space of **all possible** chess positions: 10^{40}
- Chess **cannot be solved** by supervised learning



Reinforcement Learning

Reinforcement learning (RL)

- ❑ An agent **interacts** with the world and periodically receives **rewards** (**reinforcements**) that reflect how well it is doing
 - For example, in case of chess, the reward is 1 for winning, 0 for losing, and $\frac{1}{2}$ for a draw
 - As long as we can provide the correct reward signal to the agent, **RL** provides a **very general way** to build AI systems (i.e. build **right** reward function!)
- ❑ Deep Reinforcement Learning (Deep Learning + RL) Applications



Deep Q-learning (Mnih et al., 2013)



Multi-task Robotic RL (Chebotar et al., 2021)



Poker AI (Brown et al., 2017)

Lecture 17 Reinforcement Learning

Reinforcement Learning

- ❑ Model-based **vs.** Model-free
- ❑ Passive learning agent **vs.** Active learning agent
- ❑ Exploration **vs.** Safe exploration
 - Exploration, Exploitation, Bayesian RL
- ❑ Generalization in Reinforcement Learning
 - Function approximation, Deep RL, Reward shaping, Hierarchical RL
- ❑ Apprenticeship Learning
 - Imitation learning, Inverse reinforcement learning

Outline (Lecture 17)

17.1 Learning from Rewards	7
17.2 Passive Reinforcement Learning	12
17.3 Active Reinforcement Learning	19
17.4 Generalization in Reinforcement Learning	28
17.5 Policy Search	35
17.6 Apprenticeship and Inverse Reinforcement Learning	38
17.7 Applications of Reinforcement Learning	41
Summary	43



17.1 Learning from Rewards



17.1 Learning from Rewards (1/4)

What's wrong with Supervised Learning (SL)?



- With **supervised learning**, an agent learns by passively observing example input/output pairs provided by a “**teacher**”
 - Can we train the chess agent with supervised learning?
 - Data: examples of chess positions (each labeled with the correct move)
 - Few examples of winner: 10^8
 - The space of **all possible** chess positions: 10^{40}
- Chess **cannot be solved** by **supervised learning**



17.1 Learning from Rewards (2/4)

Reinforcement Learning (RL)

- An agent interacts with the world and periodically receives **rewards** (**reinforcements**) that reflect how well it is doing
 - For example, in case of chess, the reward is 1 for winning, 0 for losing, and $\frac{1}{2}$ for a draw
 - As long as we can provide the correct reward signal to the agent, **RL** provides a **very general way** to build AI systems (i.e. build **right** reward function!)
- Deep Reinforcement Learning (Deep Learning + RL) Applications



Deep Q-learning (Mnih et al., 2013)

Multi-task Robotic RL (Chebotar et al., 2021)

Poker AI (Brown et al., 2017)

17.1 Learning from Rewards (3/4)

Categorization of Reinforcement Learning

- Model-based Reinforcement Learning
 - An agent uses a transition model of the environment.
 - The model may be **initially unknown** or **already be known**
e.g. A chess program may know the rules of chess even if it does not know how to choose good move.
 - In partially observable environments, the transition model is also useful for **state estimation**.
 - Model-based RL systems often **learn a utility function $U(s)$** , defined in terms of the sum of rewards from state s onward.

17.1 Learning from Rewards (4/4)

Categorization of Reinforcement Learning

➤ Model-free Reinforcement Learning

- An agent **neither** knows **nor** learns a transition model for the environment.
- Instead, the agent learns a more **direct representation** of how to behave.
- **Action-utility learning**
 - Most common form is **Q-learning**, where the agent learns a Q-function $Q(s, a)$
 - Given Q-function (sum of rewards from s and a), the agent can choose what to do in state s by finding the action with **the highest Q-value**
- **Policy search**
 - The agent learns a policy $\pi(s)$ that maps directly from states to actions.
 - This is a **reflex agent**.



17.2 Passive Reinforcement Learning



17.2 Passive Reinforcement Learning (1/6)

Passive learning agent


- An agent already has a **fixed** policy $\pi(s)$ that determines its actions.
- The agent is trying to learn the **utility function** $U^\pi(s)$ —the expected total discounted reward if policy π is executed beginning in state s .
- Passive learning task is similar to the **policy evaluation** task (part of **policy iteration algorithm**).
- What are the **differences** between passive learning and policy evaluation?
 - Passive learning agent does **not** know the transition model $P(s'|s, a)$.
 - Passive learning agent does **not** know the reward function $R(s, a, s')$.

17.2 Passive Reinforcement Learning (2/6)

Passive learning agent

- The optimal policies for the 4×3 world
 - $R(s, a, s') = -0.04$ (nonterminal state)

$(1,1) \xrightarrow{Up}^{-.04} (1,2) \xrightarrow{Up}^{-.04} (1,3) \xrightarrow{Right}^{-.04} (1,2) \xrightarrow{Up}^{-.04} (1,3) \xrightarrow{Right}^{-.04} (2,3) \xrightarrow{Right}^{-.04} (3,3) \xrightarrow{Right}^{+.1} (4,3)$
 $(1,1) \xrightarrow{Up}^{-.04} (1,2) \xrightarrow{Up}^{-.04} (1,3) \xrightarrow{Right}^{-.04} (2,3) \xrightarrow{Right}^{-.04} (3,3) \xrightarrow{Right}^{-.04} (3,2) \xrightarrow{Up}^{-.04} (3,3) \xrightarrow{Right}^{+.1} (4,3)$
 $(1,1) \xrightarrow{Up}^{-.04} (1,2) \xrightarrow{Up}^{-.04} (1,3) \xrightarrow{Right}^{-.04} (2,3) \xrightarrow{Right}^{-.04} (3,3) \xrightarrow{Right}^{-.04} (3,2) \xrightarrow{Up}^{-.1} (4,2)$



3	0.8516	0.9078	0.9578	<div>+1</div>
2	0.8016		0.7003	<div>-1</div>
1	0.7453	0.6953	0.6514	0.4279
	1	2	3	4

$$U^\pi(s) = E \left[\sum_{t=1}^{\infty} \gamma^t R(S_t, \pi(S_t), S_{t+1}) \right]$$

17.2 Passive Reinforcement Learning (3/6)

Direct utility estimation

- The utility of a state is defined as **the expected total reward** from that state onward (called the expected **reward-to-go**)
 - In the limit of infinitely many trials, the sample average will converge to the true expectation in equation below

$$U^\pi(s) = E \left[\sum_{t=1}^{\infty} \gamma^t R(S_t, \pi(S_t), S_{t+1}) \right]$$

- We have reduced **reinforcement learning** to a **standard supervised learning** problem in which each example is a (*state, reward-to-go*)
- But the utility of a state is determined by the reward and the expected utility of the successor states. Specifically, the utility values obey the **Bellman equations for a fixed policy**

$$U_i(s) = \sum_{s'} P(s'|s, \pi_i(s)) [R(s, \pi_i(s), s') + \gamma U_i(s')]$$

17.2 Passive Reinforcement Learning (4/6)




Adaptive dynamic programming (ADP)

- ADP agent takes advantage of the constraints among the utilities of the states by **learning the transition model that connects them and** solving the corresponding Markov decision process (MDP) using **dynamic programming**.

function PASSIVE-ADP-LEARNER(*percept*) **returns** an action
inputs: *percept*, a percept indicating the current state s' and reward signal r
persistent: π , a fixed policy
 mdp, an MDP with model P , rewards R , actions A , discount γ
 U , a table of utilities for states, initially empty
 $N_{s'|s,a}$, a table of outcome count vectors indexed by state and action, initially zero
 s, a , the previous state and action, initially null

if s' is new **then** $U[s'] \leftarrow 0$
if s is not null **then**
 increment $N_{s'|s,a}[s, a][s']$
 $R[s, a, s'] \leftarrow r$
 add a to $A[s]$
 $P(\cdot | s, a) \leftarrow \text{NORMALIZE}(N_{s'|s,a}[s, a])$
 $U \leftarrow \text{POLICY-EVALUATION}(\pi, U, \text{mdp})$
 $s, a \leftarrow s', \pi[s']$
return a

Figure 23.2 A passive reinforcement learning agent based on adaptive dynamic programming. The agent chooses a value for γ and then incrementally computes the P and R values of the MDP. The POLICY-EVALUATION function solves the fixed-policy Bellman equations, as described on page 567.



17.2 Passive Reinforcement Learning (5/6)


Temporal-difference learning (TD learning)

- Another way for solving MDP is to use the observed transitions to **adjust the utilities** of the **observed states** so that they agree with the constraint equations
 - For example, if the transition from (1,3) to (2,3): $U^\pi(1,3) = -0.04 + U^\pi(2,3)$
 - More generally: $U^\pi(s) \leftarrow U^\pi(s) + \alpha[R(s, \pi(s), s') + \gamma U^\pi(s') - U^\pi(s)]$

function PASSIVE-TD-LEARNER(*percept*) **returns** an action
inputs: *percept*, a percept indicating the current state s' and reward signal r
persistent: π , a fixed policy
 s , the previous state, initially null
 U , a table of utilities for states, initially empty
 N_s , a table of frequencies for states, initially zero


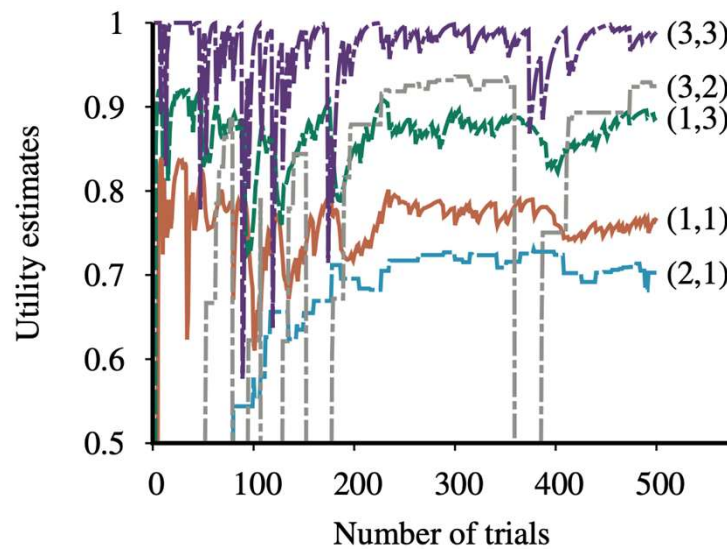
if s' is new **then** $U[s'] \leftarrow 0$
if s is not null **then**
 increment $N_s[s]$
 $U[s] \leftarrow U[s] + \alpha(N_s[s]) \times (r + \gamma U[s'] - U[s])$
 $s \leftarrow s'$
return $\pi[s']$

Figure 23.4 A passive reinforcement learning agent that learns utility estimates using temporal differences. The step-size function $\alpha(n)$ is chosen to ensure convergence.



17.2 Passive Reinforcement Learning (6/6)

Temporal-difference learning (TD learning)





17.3 Active Reinforcement Learning



17.3 Active Reinforcement Learning (1/8)

Active learning agent


- A **passive learning agent** has a **fixed policy** that determines its behavior.
- An **active learning agent** gets to decide what actions to take.
- Consider ADP agent that how it can be modified to take advantage
 - The agent will need to learn a complete transition model with outcome probabilities for all actions.
 - We need to take into account the fact that the agent has a choice of actions.
 - The utilities it needs to learn are those defined by the optimal policy; they obey the **Bellman equations** (**value iteration** or **policy iteration**)

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma U(s')]$$


17.3 Active Reinforcement Learning (2/8)

Active learning agent

- The **greedy agent** does not learn the true utilities or the true optimal policy!
Then, what should it do?



(a)




(b)

17.3 Active Reinforcement Learning (3/8)


Exploration

- The greedy agent **overlooks** the fact that actions do more than provide *reward*.
- An agent must make a **tradeoff** between *exploitation* and *exploration*.
- GLIE (Greedy in the Limit of Infinite Exploration)
 - Choose random action at time step t with probability $\frac{1}{t}$ (*exploration*)
 - Otherwise, follow the greedy policy (*exploitation*)



17.3 Active Reinforcement Learning (4/8)

Exploration

- Can we do **better** exploration?
- **A better approach:** give some **weight** to actions that the agent has **not** tried very often, while tending to avoid actions that are believed to be of **low** utility
 - $U^+(s)$: **optimistic** estimate of the utility (i.e. the expected reward-to-go)
 - $N(s, a)$: the number of times action a in state s
 - $f(u, n) = \begin{cases} R^+ & \text{if } n < N_e \\ u & \text{otherwise,} \end{cases}$  R^+ is best possible reward (optimistic estimate)
 N_e is a fixed parameter
Determine how **greed** is traded off against **curiosity**
 - Then we can rewrite the update equation with **exploration function** f as below

$$U^+(s) \leftarrow \max_a f \left(\sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma U^+(s')] , N(s, a) \right)$$

17.3 Active Reinforcement Learning (5/8)

Exploration

- An agent with exploration shows a rapid convergence toward zero policy loss (unlikely with the greedy approach)

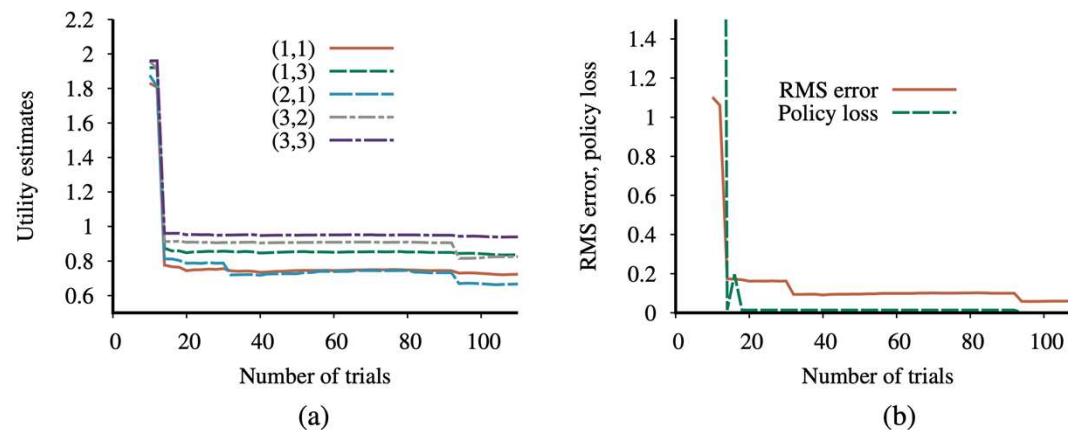


Figure 23.7 Performance of the exploratory ADP agent using $R^+ = 2$ and $N_e = 5$. (a) Utility estimates for selected states over time. (b) The RMS error in utility values and the associated policy loss.

17.3 Active Reinforcement Learning (6/8)

Safe exploration

- In a simulation environment, any accidents give us more information, and we can just hit the reset button.
- How about a real world setting? (e.g. robot learning)
 - In the worst case, the agent enters an **absorbing state** where no actions have any effect and no rewards are received.
- One of **safe exploration** approaches: **Bayesian reinforcement learning**
 - Assume a prior probability $P(h)$ over hypotheses h about the true model
 - The posterior $P(h|\mathbf{e})$ is obtained by Bayes' rule (i.e. $P(h|\mathbf{e}) = \frac{P(h)P(\mathbf{e}|h)}{P(\mathbf{e})}$)
 - The utility obtained by executing policy π in model h . Then we have,

$$\pi^* = \operatorname{argmax}_{\pi} \sum_h P(h|\mathbf{e}) U_h^{\pi}$$

17.3 Active Reinforcement Learning (7/8)

Temporal-difference Q-learning

- Two algorithms of TD learning for active ADP agent: SARSA, Q-learning
 - TD learning does **not** need a transition model $P(s'|s, a)$
- **SARSA** (for **S**tate, **A**ction, **R**eward, **S**tate, **A**ction)
 - SARSA updates with the Q-value of the action a' **that is actually taken**:
$$Q(s, a) \leftarrow Q(s, a) + \alpha [R(s, a, s') + \gamma Q(s', a') - Q(s, a)]$$
 - SARSA is an **on-policy** algorithm (behavior policy = target policy)
- **Q-learning**
 - Q-learning TD update is calculated whenever a is executed in s learning to s'
$$Q(s, a) \leftarrow Q(s, a) + \alpha [R(s, a, s') + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$
 - Q-learning is an **off-policy** algorithm (**not always** behavior policy = target policy)

17.3 Active Reinforcement Learning (8/8)

Temporal-difference Q-learning

```
function Q-LEARNING-AGENT(percept) returns an action
  inputs: percept, a percept indicating the current state  $s'$  and reward signal  $r$ 
  persistent:  $Q$ , a table of action values indexed by state and action, initially zero
                $N_{sa}$ , a table of frequencies for state–action pairs, initially zero
                $s, a$ , the previous state and action, initially null


  if  $s$  is not null then
    increment  $N_{sa}[s, a]$ 
     $Q[s, a] \leftarrow Q[s, a] + \alpha(N_{sa}[s, a])(r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$ 
     $s, a \leftarrow s', \text{argmax}_{a'} f(Q[s', a'], N_{sa}[s', a'])$   The exploration function  $f(u, n)$  we have learned
  return  $a$ 
```

Figure 23.8 An exploratory Q-learning agent. It is an active learner that learns the value $Q(s, a)$ of each action in each situation. It uses the same exploration function f as the exploratory ADP agent, but avoids having to learn the transition model.



17.4 Generalization in Reinforcement Learning



17.4 Generalization in Reinforcement Learning (1/6)

Function approximation

- Two dimensional grid environment: about 10^6 states space
 - What about **real-world** environment?
 - Backgammon is simpler than most real-world applications, yet it has 10^{20} states.
 - We **cannot** easily visit them all in order to learn how to play the game.
 - We need to use **function approximation** for approximating utility function
 - e.g. weighted linear combination of features
- $$\hat{U}_{\theta}(s) = \theta_1 f_1(s) + \theta_2 f_2(s) + \dots + \theta_n f_n(s)$$
- Instead of learning 10^6 states values (in table),
an agent can learn 20 values for parameters $\theta = (\theta_1, \theta_2, \dots, \theta_{20})$



17.4 Generalization in Reinforcement Learning (2/6)

Approximating direct utility estimation

- Consider 4×3 world environment (with x and y coordinates)

$$\hat{U}_\theta(x, y) = \theta_0 + \theta_1 x + \theta_2 y$$

- As with neural network learning, we write an **error function** and **compute its gradient** w.r.t. parameters ($u_j(s)$ is the observed total reward in j th trial)

$$E_j(s) = \left(\hat{U}_\theta(s) - u_j(s) \right)^2 / 2$$


$$\theta_i \leftarrow \theta_i - \alpha \frac{\partial E_j(s)}{\partial \theta_i} = \theta_i + \alpha [u_j(s) - \hat{U}_\theta(s)] \frac{\partial \hat{U}_\theta(s)}{\partial \theta_i}$$

- This is called the **Widrow–Hoff rule** (or **delta rule**). We get three simple update rules

$$\theta_0 \leftarrow \theta_0 + \alpha [u_j(s) - \hat{U}_\theta(s)],$$

$$\theta_1 \leftarrow \theta_1 + \alpha [u_j(s) - \hat{U}_\theta(s)] x,$$

$$\theta_2 \leftarrow \theta_2 + \alpha [u_j(s) - \hat{U}_\theta(s)] y.$$



17.4 Generalization in Reinforcement Learning (3/6)

Approximating temporal-difference learning

- We can apply the idea in previous slide equally well to TD learners

- For utility

$$\theta_i \leftarrow \theta_i + \alpha [R(s, a, s') + \gamma \hat{U}_\theta(s') - \hat{U}_\theta(s)] \frac{\partial \hat{U}_\theta(s)}{\partial \theta_i}$$

- For Q-value

$$\theta_i \leftarrow \theta_i + \alpha \left[R(s, a, s') + \gamma \max_{a'} \hat{Q}_\theta(s', a') - \hat{Q}_\theta(s, a) \right] \frac{\partial \hat{Q}_\theta(s, a)}{\partial \theta_i}$$

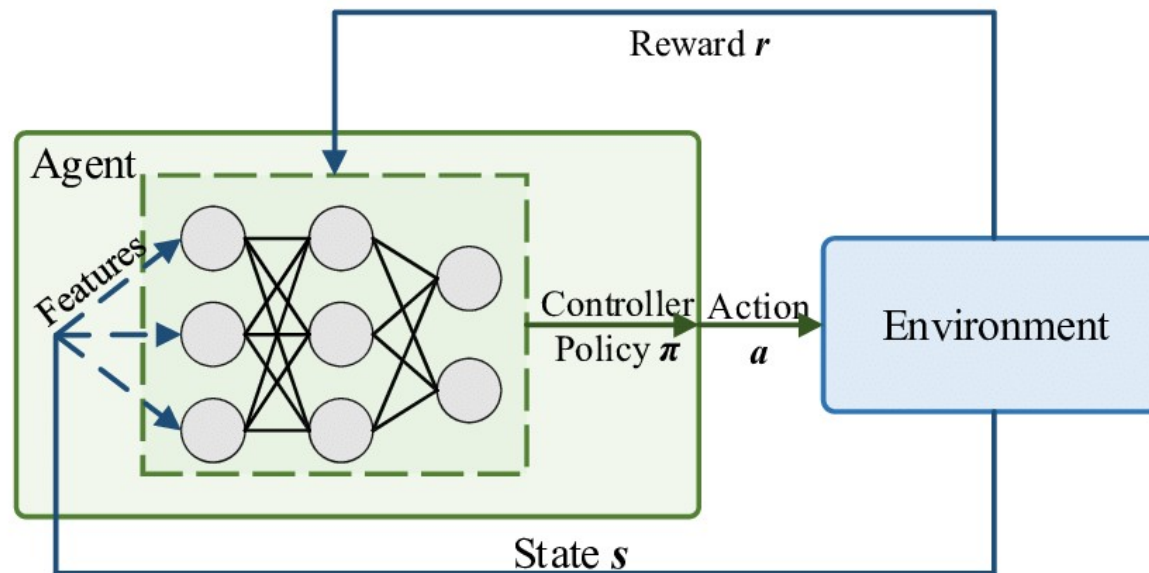
- The problem of **catastrophic forgetting**

- One of solutions is using **experience replay**.
- The learning algorithm can **retain trajectories** from the entire learning process and **replay those trajectories** to ensure that its value function is still accurate for parts of the state space it no longer visits.

17.4 Generalization in Reinforcement Learning (4/6)

Deep reinforcement learning

- The linear approximator may be insufficient → Use **deep neural networks**!
 - e.g. Video games, AlphaGo, training robots etc.




<출처> Gu et al., Machine Learning for Intelligent Optical Networks: A Comprehensive Survey

17.4 Generalization in Reinforcement Learning (5/6)

Reward shaping

- The **credit assignment** problem
 - Real-world environments may have very **sparse** rewards → many primitive actions are required to achieve any nonzero reward
 - For example, a soccer-playing robot might send a hundred thousand motor control commands to its various joints before conceding a goal. → Now it has to work out **what it did wrong**.
- One of good solutions is **reward shaping**
 - For any **potential function** Φ ,
$$R'(s, a, s') = R(s, a, s') + \gamma\Phi(s') - \Phi(s)$$
 - e.g. potential function Φ for soccer playing robot
→ A bonus for reducing distance of the ball from the opponents' goal




17.4 Generalization in Reinforcement Learning (6/6)

Hierarchical reinforcement learning

- Break **long action sequences** up into **a few smaller pieces**

```
while not IS-TERMINAL( $s$ ) do
  if BALL-IN-MY-POSSESSION( $s$ ) then choose({PASS, HOLD, DRIBBLE})
  else choose({STAY, MOVE, INTERCEPT-BALL}).
```



- HRL agent is solving a Markovian decision problem with following elements:
 - The **states** are the choice states σ of the joint state space
 - The **actions** at σ are the choices c available in σ according to the partial program
 - The **reward function** $\rho(\sigma, c, \sigma')$ is the expected sum of rewards
 - The **transition model** $\tau(\sigma, c, \sigma')$ is defined in the obvious way:
 - if c invokes a physical action a , then τ borrows from the physical model $P(s'|s, a)$;
 - if c invokes a computational transition, such as calling a subroutine, then the transition deterministically modifies the computational state m according to the rules of the programming language



17.5 Policy Search



17.5 Policy Search (1/2)

Policy representation

- Policy representation in terms of **Q-functions**

$$\pi(s) = \operatorname{argmax}_a \hat{Q}_\theta(s, a)$$

- **Problem:** policy change **discontinuously**, which makes gradient-based search difficult

- (**Stochastic**) Policy representation with **softmax**

$$\pi_\theta(s, a) = \frac{e^{\beta \hat{Q}_\theta(s, a)}}{\sum_{a'} e^{\beta \hat{Q}_\theta(s, a)}}$$

- Use a stochastic policy representation, which specifies the **probability** of actions
- Parameter $\beta > 0$ modulates softness of the softmax (**high:** hard max, **low:** uniform)

17.5 Policy Search (2/2)

Policy gradient methods

- $\rho(\theta)$ is **policy value** and $\nabla_{\theta}\rho(\theta)$ is **policy gradient**

$$\nabla_{\theta}\rho(\theta) = \nabla_{\theta} \sum_a R(s_0, a, s_0) \pi_{\theta}(s_0, a) = \sum_a R(s_0, a, s_0) \nabla_{\theta} \pi_{\theta}(s_0, a)$$

- **Monte Carlo approximation:** approximate by samples generated from $\pi_{\theta}(s_0, a)$

$$\begin{aligned} \nabla_{\theta}\rho(\theta) &= \sum_a \pi_{\theta}(s_0, a) \cdot \frac{R(s_0, a, s_0) \nabla_{\theta} \pi_{\theta}(s_0, a)}{\pi_{\theta}(s_0, a)} \\ &\approx \frac{1}{N} \sum_{j=1}^N \frac{R(s_0, a_j, s_0) \nabla_{\theta} \pi_{\theta}(s_0, a_j)}{\pi_{\theta}(s_0, a_j)} \end{aligned}$$

- For the sequential case

$$\nabla_{\theta}\rho(\theta) \approx \frac{1}{N} \sum_{j=1}^N \frac{u_j(s) \nabla_{\theta} \pi_{\theta}(s_0, a_j)}{\pi_{\theta}(s_0, a_j)}$$




17.6 Apprenticeship and Inverse Reinforcement Learning



17.6 Apprenticeship and Inverse RL (1/2)

Apprenticeship learning

- How to behave well given observations of **expert behavior**
- We can apply supervised learning using state-action pairs to learn policy
 - ➔ **Imitation learning**



<출처> Bojarski et al. '16, NVIDIA

17.6 Apprenticeship and Inverse RL (2/2)

Inverse reinforcement learning

- **Learn rewards** by observing a policy, rather than **learning a policy** by observing rewards
- How to find the **expert's reward function**, given the expert's actions
- **Feature matching** method
 - Assume the reward function as **a weighted linear combination** of features

$$R_{\theta}(s, a, s') = \sum_{i=1}^n \theta_i f_i(s, a, s') = \theta \cdot \mathbf{f}$$

- Recall the utility function of executing a policy π . We can derive **feature expectation** $\mu_i(\pi)$ (expected discounted value of the feature f_i)

$$\begin{aligned} U^{\pi}(s) &= E \left[\sum_{t=0}^{\infty} \gamma^t R(S_t, \pi(S_t), S_{t+1}) \right] = E \left[\sum_{t=0}^{\infty} \gamma^t \sum_{i=1}^n \theta_i f_i(S_t, \pi(S_t), S_{t+1}) \right] \\ &= \sum_{i=1}^n \theta_i E \left[\sum_{t=0}^{\infty} \gamma^t f_i(S_t, \pi(S_t), S_{t+1}) \right] = \sum_{i=1}^n \theta_i \mu_i(\pi) = \theta \cdot \mu(\pi) \end{aligned}$$



17.7 Applications of Reinforcement Learning



17.7 Applications of Reinforcement Learning (1/1)


Application of Reinforcement Learning

➤ Game playing

- Checker program written by Arthur Samuel (1959, 1967)
- TD-GAMMON (Backgammon) program by Gerry Tesauro (1992)
- AlphaGo program by Google DeepMind (2016)

➤ Robot control

- Cart-pole balancing problem by Michie and Chambers (1968)
- Helicopter flight using policy search by Bagnell et al. (2001)
- Deep RL for robotics and self-driving cars (2016~present)



Cart-pole balancing



Helicopter flight

Summary

1. **Direct utility estimation** uses the total observed reward-to-go for a given state as direct evidence for learning its utility.
2. A **model-based** reinforcement learning agent acquires a transition model for the environment and learns a utility function.
3. A **model-free** reinforcement learning agent may learn an action-utility function or a policy
4. **Adaptive dynamic programming** (ADP) learns a model and a reward function from observations and then uses value or policy iteration to obtain the utilities or an optimal policy.
5. **Temporal difference** (TD) methods update utility estimates to match those of successor states.
6. **Action-utility functions** (Q-functions) can be learned by an ADP approach or a TD approach.
7. **Policy-search methods** operate directly on a representation of the policy, attempting to improve it based on observed performance.
8. **Apprenticeship learning** through observation of expert behavior can be an effective solution when a correct reward function is hard to specify.
9. **Imitation learning** formulates the problem as supervised learning of a policy from the expert's state-action pairs.
10. **Inverse reinforcement learning** infers reward information from the expert's behavior.